

Heyang Sun

hsunbi@connect.ust.hk | +852 53441080 | +86 15935723459

EDUCATION

- The Hong Kong University of Science and Technology** Aug 2025 – Jun 2027
 - ✧ **MPhil in Computer Science and Engineering, Full-Scholarship Funded Research Program**
 - ✧ **Research Field:** Generative AI (LLM, MLLM, Diffusion) System **GPA:** 4.15/4.3 **Supervisor:** Wei Wang
 - ✧ **PhD Course:** Computer Vision (A+), AI system for Foundation Model (A+), Machine Learning (A), Database (A)
- Southeast University - School of Computer Science and Engineering - Bachelor** Aug 2018 – Jun 2022
 - ✧ **Major:** Software Engineering **Rank:** 2/87 **AVG:** 92.59/100 **GPA:** 3.94/4 **IELTS:** 7.5 **CET-6:** 614

WORK EXPERIENCE

- Intel - AI Frameworks Engineer (Full-time Employee, 3 Years)** Aug 2022 – Jul 2025, Shanghai
 - ✧ [IPEX LLM](#): Intel's LLM acceleration library on SGLang, Deepspeed and quantization with 8.9k stars on github
 - ✧ [BigDL](#): Distributed, Fast and Efficient AI, DL, ML and Big Data with 2.7k stars on github
 - ✧ [BigDL PPML](#): Privacy-preserving machine learning powered by Intel SGX and TDX on Kubernetes, with customers e.g. Azure, Ant, Bytedance, Tencent, AliCloud, JD and Meituan in their real-world business
- Huawei HK Research Center - Celia (小艺) AI Lab - Research Intern** Oct 2025 – Jun 2026, Hong Kong
 - ✧ [AURA](#): Built a real-time streaming video understanding framework, managing KV-cache over unbounded video with context compression, sliding window, optimized prefix caching and request scheduling
 - ✧ [BOOGU-Image](#): Built large-scale evaluation service for DiT training with vLLM, diffusers and fastAPI, and accelerations to speed up inference with cache DiT, torch compile and FP8
- Alibaba Group - Foundation Platform (爱橙) Tech - Research Intern** Sep 2025 - Now, Remote
 - ✧ Built a hybrid attention collaborating AMX and GPU to overlap KV-cache transfer for long-context LLM serving
 - ✧ Consultant on LLM accelerations, Trusted Execution Environment and other Intel and deep learning ecosystems
- TeleAI - Multi-modal Dept - World Model Inference Research Intern** Jul 2026 - Now, Shanghai
 - ✧ Improving [TeleFuser](#)'s inference performance towards models e.g. Cosmos3, Qwen-Image and Wan-Video
- Intel - Deep Learning Engineer Intern** Sep 2021 - Jul 2022, Shanghai

PUBLICATION

- ✧ VABUS: Edge-Cloud Real-time Video Analytics via Background Understanding and Subtraction. IEEE JSAC (CCF-A Journal) 2022. Second student author.
- ✧ AURA: Always-On Understanding and Real-Time Assistance via Video Streams. 2026. arXiv:2604.04184. Coauthor. **Thinking Machines Lab** cites the work as '**the closest to ours**' in their Interaction Models tech blog.
- ✧ Boogu-Image-0.1: Boosting Open-Source Unified Multimodal Understanding and Generation. 2026. arXiv:2607.13125. Coauthor.

RECOGNIZATION

- ✧ BigDL PPML: Best Customer Serving of Analytics & AI. From **Wei Li, GM & VP of Intel AIA**, 2022
- ✧ Great Effort in Helping Support **ByteDance** PPML Collaboration. From Senior Architect of Intel DCAI, 2022
- ✧ **150% Individual Performance** in 2022 Evaluation. From Manager of Intel DCAI, 2023
- ✧ **Exceeds Expectation** in 2023 Individual Performance Evaluation. From Manager of Intel DCAI, 2024
- ✧ Dedication, Technical Expertise and Professional Support to **Ant Group**. From Manager of Ant Group, 2024
- ✧ One Intel: Help **Ericsson** Fine-Tune Gemma Model at Red Hat Summit. From Manager of Intel NEX, 2025

HONOR

- ✧ President Scholarship of Southeast University (**1% and The Highest Honor in Southeast University**), Oct 2021
- ✧ Outstanding Graduate of Southeast University (**5% in SEU**), Jun 2022
- ✧ 1st Prize (**7% Nationwide**), 11th National College Mathematical Contest of China, Oct 2019
- ✧ 1st Prize (**5% in Jiangsu Province**), Advanced Mathematical Contest of Jiangsu Province, China, Jun 2019
- ✧ Excellent Graduation Thesis/Project of Jiangsu Province (**Top1 in School of Software, SEU**), Jun 2022
- ✧ Bosch Corporation AIoT Scholarship, 2021; HUAWEI Scholarship, 2021; NUODE Logistics Scholarship, 2020